

## Exercise 1: PRACTICE PANDAS LIBRARY ON diabetes DATASET

### Importing pandas

To begin working with pandas, import the pandas Python package as shown below. When importing pandas, the most common alias for pandas is pd.

```
import pandas as pd
```

### Load Dataset

### Mounting Google Drive

```
from google.colab import drive
drive.mount('/content/drive')
```

```
-----
ModuleNotFoundError                                Traceback (most recent call last)
c:\Users\aadit\Desktop\BTech\S5\Foundations Of Data Science\LAB\lab 1\Untitled.ipynb
Cell 6 line 1
----> <a href='vscode-notebook-cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20Of%20Data%20Science/LAB/lab%201/Ur
line=0'>1</a> from google.colab import drive
      <a href='vscode-notebook-cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20Of%20Data%20Science/LAB/lab%201/Ur
line=1'>2</a> drive.mount('/content/drive')
```

```
ModuleNotFoundError: No module named 'google'
```

### Importing csv files

Use read\_csv() with the path to the CSV file to read a comma-separated values file.

```
df=pd.read_csv('diabetes.csv')
```

### Importing Text Files

Reading text files is similar to CSV files. The only nuance is that you need to specify a separator with the `sep` argument, as shown below.

The separator argument refers to the symbol used to separate rows in a Df.

Comma (`sep = ","`), whitespace (`sep = "\s"`), tab (`sep = "\t"`), and colon (`sep = ":"`) are the commonly used separators. Here `\s` represents a single white space character.

```
df=pd.read_csv('diabetes.txt',sep='\s')
```

### ▼ Importing Excel files (single sheet)

Reading excel files (both XLS and XLSX) is as easy as the `read_excel()` function, using the file path as an input.

```
df=pd.read_excel('diabetes.xlsx')
```

### ▼ Viewing and understanding Dfs using pandas

After reading tabular data as a Df, you would need to have a glimpse of the data. You can either view a small sample of the dataset or a summary of the data in the form of summary statistics.

### ▼ How to view data using `.head()` and `.tail()`

You can view the first few or last few rows of a Df using the `.head()` or `.tail()` methods, respectively. You can specify the number of rows through the `n` argument (the default value is 5).

```
df.head()
```

Pregnancies Glucose BloodPressure SkinThickness Insulin BMI DiabetesPedigreeFu

```
df.tail(n=9)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
<b>759</b>	6	190	92	0	0	35.5	
<b>760</b>	2	88	58	26	16	28.4	
<b>761</b>	9	170	74	31	0	44.0	
<b>762</b>	9	89	62	0	0	22.5	
<b>763</b>	10	101	76	48	180	32.9	
<b>764</b>	2	122	70	27	0	36.8	
<b>765</b>	5	121	72	23	112	26.2	
<b>766</b>	1	126	60	0	0	30.1	
<b>767</b>	1	93	70	31	0	30.4	

## ▼ Understanding data using .describe()

The .describe() method prints the summary statistics of all numeric columns, such as count, mean, standard deviation, range, and quartiles of numeric columns.

```
df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Dia
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

You can also modify the quartiles using the percentiles argument. Here, for example, we're looking at the 30%, 50%, and 70% percentiles of the numeric columns in Df df

```
df.describe(percentiles=[0.3,0.5,0.7])
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Dia
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
<b>30%</b>	1.000000	102.000000	64.000000	8.200000	0.000000	28.200000	
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
<b>70%</b>	5.000000	134.000000	78.000000	31.000000	106.000000	35.490000	
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

You can isolate specific data types in your summary output by using the include argument. Here, for example, we're only summarizing the columns with the integer data type.

```
df.describe(include=[int])
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	Age	(
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	33.240885	0
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	11.760232	0
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	21.000000	0
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	24.000000	0
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	29.000000	0
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	41.000000	1
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	81.000000	1

Try transposing them with the .T attribute.

```
df.describe().T
```

	count	mean	std	min	25%	50%	75%
<b>Pregnancies</b>	768.0	3.845052	3.369578	0.000	1.00000	3.0000	6.0000
<b>Glucose</b>	768.0	120.894531	31.972618	0.000	99.00000	117.0000	140.2500
<b>BloodPressure</b>	768.0	69.105469	19.355807	0.000	62.00000	72.0000	80.0000
<b>SkinThickness</b>	768.0	20.536458	15.952218	0.000	0.00000	23.0000	32.0000
<b>Insulin</b>	768.0	79.799479	115.244002	0.000	0.00000	30.5000	127.2500
<b>BMI</b>	768.0	31.992578	7.884160	0.000	27.30000	32.0000	36.6000
<b>DiabetesPedigreeFunction</b>	768.0	0.471876	0.331329	0.078	0.24375	0.3725	0.6250
<b>Age</b>	768.0	33.240885	11.760232	21.000	24.00000	29.0000	41.0000
<b>Outcome</b>	768.0	0.348958	0.476951	0.000	0.00000	0.0000	1.0000

## ▼ Understanding data using .info()

The .info() method is a quick way to look at the data types, missing values, and data size of a DataFrame.. When verbose is set to True, it prints the full summary from .info().

```
df.info(show_counts=True,memory_usage=True,verbose=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome               768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

## ▼ Understanding your data using .shape

The number of rows and columns of a DataFrame can be identified using the .shape attribute of the DataFrame. It returns a tuple (row, column) and can be indexed

```
df.shape # Get the number of rows and columns
```

```
(768, 9)
```

```
df.shape[0] # Get the number of rows only
```

```
768
```

```
df.shape[1] # Get the number of columns onlyd to get only rows, and only columns count as ou
```

```
9
```

## ▼ Get all columns and column names

Calling the `.columns` attribute of a DataFrame object returns the column names in the form of an Index object. As a reminder, a pandas index is the address/label of the row or column.

```
df.columns
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',  
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],  
      dtype='object')
```

## ▼ Checking for missing values in pandas with `.isnull()`

The sample DataFrame does not have any missing values. Let's introduce a few to make things interesting. The `.copy()` method makes a copy of the original DataFrame. This is done to ensure that any changes to the copy don't reflect in the original DataFrame. Using `.loc` (to be discussed later), you can set rows two to five of the `Pregnancies` column to `NaN` values, which denote missing values.

```
df2 = df.copy()
```

```
df2.loc[2:5, 'Pregnancies'] = None  
df2.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6.0	148	72	35	0	33.6	
1	1.0	85	66	29	0	26.6	

```
df2.head(7)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6.0	148	72	35	0	33.6	
1	1.0	85	66	29	0	26.6	
2	NaN	183	64	0	0	23.3	
3	NaN	89	66	23	94	28.1	
4	NaN	137	40	35	168	43.1	
5	NaN	116	74	0	0	25.6	
6	3.0	78	50	32	88	31.0	

You can check whether each element in a DataFrame is missing using the `.isnull()` method.

```
df2.isnull().head(7)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	True	False	False	False	False	False	
3	True	False	False	False	False	False	
4	True	False	False	False	False	False	
5	True	False	False	False	False	False	
6	False	False	False	False	False	False	

Given it's often more useful to know how much missing data you have, you can combine `.isnull()` with `.sum()` to count the number of nulls in each column.

```
df2.isnull().sum()
```

```
Pregnancies    4
Glucose        0
```

```

BloodPressure      0
SkinThickness      0
Insulin            0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64

```

You can also do a double sum to get the total number of nulls in the DataFrame.

```
df2.isnull().sum().sum()
```

```
4
```

## ▼ Slicing and Extracting Data in pandas

The pandas package offers several ways to subset, filter, and isolate data in your DataFrames. Here, we'll see the most common ways.

### ▼ • Isolating one column using [ ]

You can isolate a single column using a square bracket [ ] with a column name in it. The output is a pandas Series object. A pandas Series is a one-dimensional array containing data of any type, including integer, float, string, boolean, python objects, etc. A DataFrame is comprised of many series that act as columns.

```
df['Outcome']
```

```

0      1
1      0
2      1
3      0
4      1
..
763    0
764    0
765    0
766    1
767    0
Name: Outcome, Length: 768, dtype: int64

```

### ▼ Isolating two or more columns using [ ]



You can also provide a list of column names inside the square brackets to fetch more than one column. Here, square brackets are used in two different ways. We use the outer square brackets to indicate a subset of a DataFrame, and the inner square brackets to create a list.

```
df[['Pregnancies', 'Outcome']]
```

	Pregnancies	Outcome
0	6	1
1	1	0
2	8	1
3	1	0
4	0	1
...	...	...
763	10	0
764	2	0
765	5	0
766	1	1
767	1	0

768 rows × 2 columns

### ▼ Isolating one row using [ ]

A single row can be fetched by passing in a boolean series with one True value. In the example below, the second row with index = 1 is returned. Here, `.index` returns the row labels of the DataFrame, and the comparison turns that into a Boolean one-dimensional array.

```
df[df.index==1]
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
1	1	85	66	29	0	26.6	

### ▼ Isolating two or more rows using [ ]

Similarly, two or more rows can be returned using the `.isin()` method instead of a `==` operator.

```
df[df.index.isin(range(2,10))]
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
5	5	116	74	0	0	25.6	
6	3	78	50	32	88	31.0	
7	10	115	0	0	0	35.3	
8	2	197	70	45	543	30.5	
9	8	125	96	0	0	0.0	

## ▼ Using `.loc[]` and `.iloc[]` to fetch rows

You can fetch specific rows by labels or conditions using `.loc[]` and `.iloc[]` ("location" and "integer location"). `.loc[]` uses a label to point to a row, column or cell, whereas `.iloc[]` uses the numeric position. To understand the difference between the two, let's modify the index of `df2` created earlier.

```
df2.index = range(1,769)
```

The below example returns a pandas Series instead of a DataFrame. The 1 represents the row index (label), whereas the 1 in `.iloc[]` is the row position (first row).

```
df2.loc[1]
```

```
Pregnancies      6.000
Glucose          148.000
BloodPressure     72.000
SkinThickness     35.000
Insulin           0.000
BMI              33.600
DiabetesPedigreeFunction  0.627
Age              50.000
Outcome          1.000
Name: 1, dtype: float64
```

```
df2.iloc[1]
```

```
Pregnancies      1.000
Glucose           85.000
BloodPressure     66.000
SkinThickness     29.000
Insulin           0.000
BMI               26.600
DiabetesPedigreeFunction  0.351
Age               31.000
Outcome           0.000
Name: 2, dtype: float64
```

You can also fetch multiple rows by providing a range in square brackets.

```
df2.loc[100:110]
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
<b>100</b>	1.0	122	90	51	220	49.7	
<b>101</b>	1.0	163	72	0	0	39.0	
<b>102</b>	1.0	151	60	0	0	26.1	
<b>103</b>	0.0	125	96	0	0	22.5	
<b>104</b>	1.0	81	72	18	40	26.6	
<b>105</b>	2.0	85	65	0	0	39.6	
<b>106</b>	1.0	126	56	29	152	28.7	
<b>107</b>	1.0	96	122	0	0	22.4	
<b>108</b>	4.0	144	58	28	140	29.5	
<b>109</b>	3.0	83	58	31	18	34.3	
<b>110</b>	0.0	95	85	25	36	37.4	

You can also subset with `.loc[]` and `.iloc[]` by using a list instead of a range.

```
df2.loc[[100, 200, 300]]
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
<b>100</b>	1.0	122	90	51	220	49.7	
<b>200</b>	4.0	148	60	27	318	30.9	
<b>300</b>	8.0	112	72	0	0	23.6	

You can also select specific columns along with rows. This is where `.iloc[]` is different from `.loc[]` – it requires column location and not column labels.

```
df2.loc[100:110, ['Pregnancies', 'Glucose', 'BloodPressure']]
```

	Pregnancies	Glucose	BloodPressure
<b>100</b>	1.0	122	90
<b>101</b>	1.0	163	72
<b>102</b>	1.0	151	60
<b>103</b>	0.0	125	96
<b>104</b>	1.0	81	72
<b>105</b>	2.0	85	65
<b>106</b>	1.0	126	56
<b>107</b>	1.0	96	122
<b>108</b>	4.0	144	58
<b>109</b>	3.0	83	58
<b>110</b>	0.0	95	85

For faster workflows, you can pass in the starting index of a row as a range.

```
df2.loc[760:, ['Pregnancies', 'Glucose', 'BloodPressure']]
```

	Pregnancies	Glucose	BloodPressure
<b>760</b>	6.0	190	92
<b>761</b>	2.0	88	58
<b>762</b>	9.0	170	74
<b>763</b>	9.0	89	62
<b>764</b>	10.0	101	76
<b>765</b>	2.0	122	70
<b>766</b>	5.0	121	72
<b>767</b>	1.0	126	60
<b>768</b>	1.0	93	70

## ▼ Conditional slicing (that fits certain conditions)

pandas lets you filter data by conditions over row/column values. For example, the below code selects the row where Blood Pressure is exactly 122. Here, we are isolating rows using the brackets `[]` as seen in previous sections. However, instead of inputting row indices or column names, we are inputting a condition where the column BloodPressure is equal to 122. We denote this condition using `df.BloodPressure == 122`.

```
df[df.BloodPressure == 122]
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
<b>106</b>	1	96	122	0	0	22.4	

The below example fetched all rows where Outcome is 1. Here `df.Outcome` selects that column, `df.Outcome == 1` returns a Series of Boolean values determining which Outcomes are equal to 1, then `[]` takes a subset of `df` where that Boolean Series is True.

```
df[df.Outcome == 1]
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
<b>0</b>	6	148	72	35	0	33.6	
<b>2</b>	8	183	64	0	0	23.3	
<b>4</b>	0	137	40	35	168	43.1	
<b>6</b>	3	78	50	32	88	31.0	
<b>8</b>	2	197	70	45	543	30.5	
...	...	...	...	...	...	...	
<b>755</b>	1	128	88	39	110	36.5	
<b>757</b>	0	123	72	0	0	36.3	
<b>759</b>	6	190	92	0	0	35.5	
<b>761</b>	9	170	74	31	0	44.0	
<b>766</b>	1	126	60	0	0	30.1	

268 rows × 9 columns

You can use a `>` operator to draw comparisons. The below code fetches Pregnancies, Glucose, and BloodPressure for all records with BloodPressure greater than 100.

```
df.loc[df['BloodPressure'] > 100, ['Pregnancies', 'Glucose', 'BloodPressure' ]]
```

	Pregnancies	Glucose	BloodPressure
<b>43</b>	9	171	110
<b>84</b>	5	137	108
<b>106</b>	1	96	122
<b>177</b>	0	129	110
<b>207</b>	5	162	104
<b>362</b>	5	103	108
<b>369</b>	1	133	102
<b>440</b>	0	189	104
<b>549</b>	4	189	110
<b>658</b>	11	127	106
<b>662</b>	8	167	106
<b>672</b>	10	68	106
<b>691</b>	13	158	114

## ▼ Data analysis in pandas

The main value proposition of pandas lies in its quick data analysis functionality. In this section, we'll focus on a set of analysis techniques you can use in pandas.

## ▼ Summary operators (mean, mode, median)

As you saw earlier, you can get the mean of each column value using the `.mean()` method.

```
df.mean()
```

```
Pregnancies    3.845052
Glucose        120.894531
BloodPressure   69.105469
SkinThickness   20.536458
```

```

Insulin          79.799479
BMI              31.992578
DiabetesPedigreeFunction  0.471876
Age              33.240885
Outcome          0.348958
dtype: float64

```

## ▼ Create new columns based on existing columns

pandas provides fast and efficient computation by combining two or more columns like scalar variables. The below code divides each value in the column Glucose with the corresponding value in the Insulin column to compute a new column named Glucose\_Insulin\_Ratio.

```
df2['Glucose_Insulin_Ratio'] = df2['Glucose']/df2['Insulin']
```

```
df2.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
1	6.0	148	72	35	0	33.6	
2	1.0	85	66	29	0	26.6	
3	NaN	183	64	0	0	23.3	
4	NaN	89	66	23	94	28.1	
5	NaN	137	40	35	168	43.1	

## ▼ Counting using .value\_counts()

Often times you'll work with categorical values, and you'll want to count the number of observations each category has in a column. Category values can be counted using the .value\_counts() methods. Here, for example, we are counting the number of observations where Outcome is diabetic (1) and the number of observations where the Outcome is nondiabetic (0).

```
df['Outcome'].value_counts()
```

```

Outcome
0      500
1      268
Name: count, dtype: int64

```

## ▼ Aggregating data with .groupby() in pandas

pandas lets you aggregate values by grouping them by specific column values. You can do that by combining the .groupby() method with a summary method of your choice. The code below displays the mean of each numeric column grouped by Outcome.

```
df.groupby('Outcome').mean()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Di
Outcome							
0	3.298000	109.980000	68.184000	19.664000	68.792000	30.304200	
1	4.865672	141.257463	70.824627	22.164179	100.335821	35.142537	

.groupby() enables grouping by more than one column by passing a list of column names, as shown below

```
df.groupby(['Pregnancies', 'Outcome']).mean()
```



		Glucose	BloodPressure	SkinThickness	Insulin	BMI	Di
Pregnancies	Outcome						
0	0	111.945205	69.205479	21.054795	77.561644	31.727397	
	1	144.236842	63.210526	24.605263	89.578947	39.213158	
1	0	104.254717	66.830189	23.047170	84.320755	29.616038	
	1	143.793103	71.310345	29.517241	151.137931	37.793103	
2	0	105.214286	61.940476	20.107143	72.619048	29.679762	
	1	135.473684	69.052632	28.210526	144.315789	34.578947	
3	0	109.604167	65.708333	17.520833	62.020833	29.231250	
	1	148.444444	68.148148	24.629630	132.666667	32.548148	
4	0	117.555556	71.577778	18.422222	78.466667	31.255556	
	1	139.913043	67.000000	10.913043	51.782609	33.873913	
5	0	111.666667	74.666667	17.166667	46.861111	31.100000	
	1	131.190476	78.857143	17.761905	75.190476	36.780952	
6	0	115.352941	66.382353	18.705882	69.029412	29.591176	
	1	132.375000	72.750000	15.375000	52.000000	31.775000	
7	0	121.000000	70.350000	19.350000	72.500000	29.975000	
	1	148.800000	71.120000	21.040000	94.040000	34.756000	

Any summary method can be used alongside `.groupby()`, including `.min()`, `.max()`, `.mean()`, `.median()`, `.sum()`, `.mode()`, and more.

9	0	107.000000	70.400000	22.400000	71.200000	28.840000
---	---	------------	-----------	-----------	-----------	-----------

## ▼ Exercise 2: PRACTICE PANDAS EXERCISE FROM

<https://www.w3resource.com/python-exercises/pandas/index.php>

## ▼ Exercise 3: Using Wine Quality dataset do the following:

1	116.750000	71.500000	30.250000	213.500000	34.575000
---	------------	-----------	-----------	------------	-----------

1. Copy the given winequality dataset to your local folder.

1	133.000000	73.200000	12.000000	3.000000	30.720000
---	------------	-----------	-----------	----------	-----------

2. Load the winequality dataset using pandas

```
import pandas as pd
df = pd.read_csv('WineQT.csv')
```

### 3. Find the size of the dataset

```
df.shape
```

```
(1143, 13)
```

### 4. Get the statistical summary of the data

```
df.describe()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
<b>count</b>	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000
<b>mean</b>	8.311111	0.531339	0.268364	2.532152	0.086933	15.615486	45.970321
<b>std</b>	1.747595	0.179633	0.196686	1.355917	0.047267	10.250486	32.765549
<b>min</b>	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000
<b>25%</b>	7.100000	0.392500	0.090000	1.900000	0.070000	7.000000	21.000000
<b>50%</b>	7.900000	0.520000	0.250000	2.200000	0.079000	13.000000	37.000000
<b>75%</b>	9.100000	0.640000	0.420000	2.600000	0.090000	21.000000	61.000000

### 5. Print the first and last five rows

```
df.head()
```

free total

df.tail()

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	su
1138	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	
1139	6.8	0.620	0.08	1.9	0.068	28.0	38.0	0.99651	3.42	
1140	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	
1141	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	

6. Print the first 7 rows of the dataset.

df.head(7)

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulph
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	

7. Display the first 7 rows of 4th column to 6th column.

df.head(7).T.head(6).tail(3).T

	residual sugar	chlorides	free sulfur dioxide
<b>0</b>	1.9	0.076	11.0

8. Select all the rows of the 3rd and 6th column

```
df[['citric acid', 'free sulfur dioxide' ]]
```

	citric acid	free sulfur dioxide
<b>0</b>	0.00	11.0
<b>1</b>	0.00	25.0
<b>2</b>	0.04	15.0
<b>3</b>	0.56	17.0
<b>4</b>	0.00	11.0
...	...	...
<b>1138</b>	0.13	29.0
<b>1139</b>	0.08	28.0
<b>1140</b>	0.08	32.0
<b>1141</b>	0.10	39.0
<b>1142</b>	0.12	32.0

1143 rows × 2 columns

9. Check the number of null and non-null values in the dataset columns wise.

```
df.isnull().sum()
```

```
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                0
sulphates          0
alcohol            0
quality            0
Id                 0
dtype: int64
```

```
df.notnull().sum()
```

```
fixed acidity      1143
volatile acidity   1143
citric acid        1143
residual sugar     1143
chlorides          1143
free sulfur dioxide 1143
total sulfur dioxide 1143
density            1143
pH                1143
sulphates          1143
alcohol            1143
quality            1143
Id                 1143
dtype: int64
```

## 10. Rename the pH column as potential of hydrogen

```
df.rename(columns={'pH': 'potential of hydrogen'}, inplace=True)
```

## 11. Create a new column named total\_free ratio which is the ratio of the total sulphur dioxide and free sulphurdioxide

```
df['total_free ratio']=df['free sulfur dioxide']/df['total sulfur dioxide']
```

```
df.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	potential of hydrogen
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16

## 12. Aggregate the data based on quality

```
df.describe()['quality']
```

```

count      1143.000000
mean        5.657043
std         0.805824
min         3.000000
25%         5.000000
50%         6.000000
75%         6.000000
max         8.000000
Name: quality, dtype: float64

```

13. Slice the data frame and create a new one such that the rows containing total\_free ratio values less than 2.7 and greater than 3.2

```

df2 = df.copy()
df2 = df2[(df2['total_free ratio'] < 2.7) | (df2['total_free ratio'] > 3.2)]

```

14. Drop second column values from the dataset

```
df2.drop(['volatile acidity'], inplace=True, axis=1)
```

df2

	fixed acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	potential of hydrogen	sulpha
<b>0</b>	7.4	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0
<b>1</b>	7.8	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0
<b>2</b>	7.8	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0
<b>3</b>	11.2	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0
<b>4</b>	7.4	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0
...	...	...	...	...	...	...	...	...	...
<b>1138</b>	6.3	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0
<b>1139</b>	6.8	0.08	1.9	0.068	28.0	38.0	0.99651	3.42	0
<b>1140</b>	6.2	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0
<b>1141</b>	5.9	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0
<b>1142</b>	5.9	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0

15. Sort the DataFrame first by " in descending order, then by " in ascending order

16. Find the sum and the cumulative sum of second attribute

```
sum_second_attribute = df['citric acid'].sum()
cumulative_sum_second_attribute = df['citric acid'].cumsum()

print(sum_second_attribute, cumulative_sum_second_attribute)
```

```
306.74 0          0.00
1          0.00
2          0.04
3          0.60
4          0.60
...
1138      306.36
1139      306.44
1140      306.52
1141      306.62
1142      306.74
Name: citric acid, Length: 1143, dtype: float64
```

17. Find the minimum and maximum of third attribute

```
minimum = df['residual sugar'].min()
maximum = df['residual sugar'].max()

print("Minimum:", minimum)
print("Maximum:", maximum)
```

```
Minimum: 0.9
Maximum: 15.5
```

18. Rename columns of a DataFrame

```
df.rename(columns={'potential of hydrogen': 'ph'}, inplace=True)
df.columns
```

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
```

```
'ph', 'sulphates', 'alcohol', 'quality', 'Id', 'total_free ratio'],  
dtype='object')
```

19. Change the order of a DataFrame columns.

```
df = df[['Id', 'fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',  
        'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',  
        'ph', 'sulphates', 'alcohol', 'quality', 'total_free ratio']]
```

20. Write a DataFrame to CSV file using tab separator.

```
df.to_csv('output_file.csv', sep='\t', index=False)
```

21. Replace all the NaN values with Zero's in a column of a dataframe

```
df.isnull().sum()
```

```
Id          0  
fixed acidity    0  
volatile acidity    0  
citric acid      0  
residual sugar    0  
chlorides        0  
free sulfur dioxide    0  
total sulfur dioxide    0  
density          0  
ph              0  
sulphates        0  
alcohol          0  
quality          0  
total_free ratio    0  
dtype: int64
```

```
df['fixed acidity'].fillna(0, inplace=True)
```

22. Divide a data frame in 60:40 ratio.



```
from sklearn.model_selection import train_test_split

train_df, test_df = train_test_split(df, test_size=0.4, random_state=42)

print("Train set shape:", train_df.shape)
print("Test set shape:", test_df.shape)
```

```
Train set shape: (685, 14)
Test set shape: (458, 14)
```

### 23. Combine two series into a DataFrame.

```
import pandas as pd

series1 = pd.Series([1, 2, 3])
series2 = pd.Series([4, 5, 6])

df = pd.concat([series1, series2], axis=1)

print(df)
```

```
   0  1
0  1  4
1  2  5
2  3  6
```

### 24. Find the row for where the value of a first attribute is maximum.

```
max_row = df[df['residual sugar'] == df['residual sugar'].max()]
print(max_row)
```

```
fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  \
339          10.6             0.28         0.39           15.5         0.069

free sulfur dioxide  total sulfur dioxide  density    pH  sulphates  \
339                6.0                23.0   1.0026  3.12         0.66

alcohol  quality  Id
339      9.2      5  480
```

### 25. Get the datatypes of columns of a DataFrame.

```
df.dtypes
```



fixed acidity	float64
volatile acidity	float64
citric acid	float64
residual sugar	float64
chlorides	float64
free sulfur dioxide	float64
total sulfur dioxide	float64
density	float64
pH	float64
sulphates	float64
alcohol	float64
quality	int64
Id	int64
dtype: object	