

Lab Sheet 1

Exercise 1

1. Read the csv file from the given URL to a dataframe. <https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data>.

```
from pandas import read_csv
path='imports-85.data'
data=read_csv(path, )
data.head()
```

	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.60	...	130	mpfi	3.47	2
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	:
1	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	:
2	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	:
3	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	:

2. Include the following headers to the above csv file ["symboling","normalized-losses","make","fuel-type","aspiration", "num-of-doors","body-style","drive-wheels","engine-location","wheel-base","length","width","height","curb-weight","engine-type","num-of-cylinders", "engine-size","fuel-system","bore","stroke","compression-ratio","horsepower", "peak-rpm","city-mpg","highway-mpg","price"]

```
import pandas as pd
head = ["symboling","normalized-losses","make","fuel-type","aspiration", "num-of-doors","body-style","drive-wheels","engine-location","wheel-base","length","width","height","curb-weight","engine-type","num-of-cylinders", "engine-size","fuel-system","bore","stroke","compression-ratio","horsepower", "peak-rpm","city-mpg","highway-mpg","price"]
df = pd.read_csv(path, names = head)
df.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front
3	2	164	audi	gas	std	four	sedan	fwd	front
4	2	164	audi	gas	std	four	sedan	4wd	front

5 rows × 26 columns

3. Display the first five rows of the dataset.

```
df.head(5)
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front
3	2	164	audi	gas	std	four	sedan	fwd	front
4	2	164	audi	gas	std	four	sedan	4wd	front

205 rows × 10 columns

- ▼ 4. Replace the "?" in the above file with NaN.

```
import numpy as np
df.replace("?", np.nan, inplace = True)
```

- ▼ 5. Find the missing values in the dataset.

```
df.isnull()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wh
0	False	True	False	False	False	False	False	False	False	l
1	False	True	False	False	False	False	False	False	False	l
2	False	True	False	False	False	False	False	False	False	l
3	False	False	False	False	False	False	False	False	False	l
4	False	False	False	False	False	False	False	False	False	l
...
200	False	False	False	False	False	False	False	False	False	l
201	False	False	False	False	False	False	False	False	False	l
202	False	False	False	False	False	False	False	False	False	l
203	False	False	False	False	False	False	False	False	False	l
204	False	False	False	False	False	False	False	False	False	l

205 rows × 11 columns

- ▼ 6. Count the missing values in each column.

```
df.isnull().sum()
```

```
symboling      0
normalized-losses  41
make           0
fuel-type      0
aspiration     0
num-of-doors   2
body-style     0
drive-wheels   0
engine-location 0
wheel-base    0
length        0
width         0
height        0
curb-weight    0
engine-type    0
num-of-cylinders 0
engine-size    0
```

```

fuel-system      0
bore             4
stroke          4
compression-ratio 0
horsepower      2
peak-rpm        2
city-mpg        0
highway-mpg     0
price           4
dtype: int64

```

- ▼ 7. Identify the column(s) of a given Data Frame which have at least one missing value.

```
df.columns[df.isnull().any()]
```

```

Index(['normalized-losses', 'num-of-doors', 'bore', 'stroke', 'horsepower',
      'peak-rpm', 'price'],
      dtype='object')

```

- ▼ 8. Find the Indexes of missing values of column “normalized-losses” in the given DataFrame.

```
df[df["normalized-losses"].isnull()].index
```

```

Index([ 0, 1, 2, 5, 7, 9, 14, 15, 16, 17, 43, 44, 45, 46,
      48, 49, 63, 66, 71, 73, 74, 75, 82, 83, 84, 109, 110, 113,
      114, 124, 126, 127, 128, 129, 130, 131, 181, 189, 191, 192, 193],
      dtype='int64')

```

- ▼ 9. Replace the missing values in “normalized-losses”, “stroke”, “bore”, with the mean.

```

df["normalized-losses"] = pd.to_numeric(df["normalized-losses"], errors='coerce')
df["stroke"] = pd.to_numeric(df["stroke"], errors='coerce')
df["bore"] = pd.to_numeric(df["bore"], errors='coerce')

```

```

mean_normalized_losses = df["normalized-losses"].mean()
mean_stroke = df["stroke"].mean()
mean_bore = df["bore"].mean()

```

```

df["normalized-losses"].fillna(mean_normalized_losses, inplace=True)
df["stroke"].fillna(mean_stroke, inplace=True)
df["bore"].fillna(mean_bore, inplace=True)

```

```
print(df.head())
```

```

   symboling  normalized-losses      make fuel-type aspiration \
0          3             122.0  alfa-romero      gas      std
1          3             122.0  alfa-romero      gas      std
2          1             122.0  alfa-romero      gas      std
3          2             164.0      audi      gas      std
4          2             164.0      audi      gas      std

   num-of-doors  body-style  drive-wheels  engine-location  wheel-base  ... \
0          two  convertible      rwd      front      88.6  ...
1          two  convertible      rwd      front      88.6  ...
2          two   hatchback      rwd      front      94.5  ...
3         four      sedan      fwd      front      99.8  ...
4         four      sedan      4wd      front      99.4  ...

   engine-size  fuel-system  bore  stroke  compression-ratio  horsepower \
0          130      mpfi  3.47   2.68           9.0          111
1          130      mpfi  3.47   2.68           9.0          111
2          152      mpfi  2.68   3.47           9.0          154
3          109      mpfi  3.19   3.40          10.0          102
4          136      mpfi  3.19   3.40           8.0          115

   peak-rpm  city-mpg  highway-mpg  price
0       5000       21          27  13495
1       5000       21          27  16500
2       5000       19          26  16500
3       5500       24          30  13950
4       5500       18          22  17450

```

[5 rows x 26 columns]

11. Replace the missing values in "num_doors" with the maximum frequency value

```
max_freq = df['num-of-doors'].mode()[0]
df['num-of-doors'].fillna(max_freq, inplace=True)
df.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location
0	3	122.0	alfa-romero	gas	std	two	convertible	rwd	front
1	3	122.0	alfa-romero	gas	std	two	convertible	rwd	front
2	1	122.0	alfa-romero	gas	std	two	hatchback	rwd	front
3	2	164.0	audi	gas	std	four	sedan	fwd	front
4	2	164.0	audi	gas	std	four	sedan	4wd	front

5 rows x 26 columns

12. Replace the missing values in "horse_power", "peak_rpm" with the values in the next row.

```
df[['horsepower', 'peak-rpm']] = df[['horsepower', 'peak-rpm']].fillna(method='bfill')
df.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location
0	3	122.0	alfa-romero	gas	std	two	convertible	rwd	front
1	3	122.0	alfa-romero	gas	std	two	convertible	rwd	front
2	1	122.0	alfa-romero	gas	std	two	hatchback	rwd	front
3	2	164.0	audi	gas	std	four	sedan	fwd	front
4	2	164.0	audi	gas	std	four	sedan	4wd	front

5 rows x 26 columns

13. Drop the rows of "price", if value is missing.

```
df.dropna(subset=["price"], axis=0, inplace=True)
```

```
df.head()
```

symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	
0	3	122.0	alfa-	gas	std	two	convertible	rwd	front

▼ 14. List the datatypes of each column.

```
print(df.dtypes)
```

```

symboling          int64
normalized-losses  float64
make              object
fuel-type         object
aspiration        object
num-of-doors      object
body-style        object
drive-wheels      object
engine-location   object
wheel-base       float64
length           float64
width            float64
height           float64
curb-weight       int64
engine-type       object
num-of-cylinders  object
engine-size       int64
fuel-system       object
bore             float64
stroke           float64
compression-ratio float64
horsepower        object
peak-rpm          object
city-mpg          int64
highway-mpg       int64
price            object
dtype: object

```

▼ 15. Convert the columns to appropriate datatype.

```

df[["symboling", "normalized-losses"]] = df[["symboling", "normalized-losses"]].astype("int")
df[["bore", "stroke", "price", "peak-rpm"]] = df[["bore", "stroke", "price", "peak-rpm"]].astype("float")
df[["num-of-doors", "num-of-cylinders"]] = df[["num-of-doors", "num-of-cylinders"]].replace({"four": 4, "six": 6, "five": 5, "eight": 8, "two": 2})
df[["num-of-doors", "num-of-cylinders"]] = df[["num-of-doors", "num-of-cylinders"]].astype("int")
print(df.dtypes)

```

```

symboling          int32
normalized-losses  int32
make              object
fuel-type         object
aspiration        object
num-of-doors      int32
body-style        object
drive-wheels      object
engine-location   object
wheel-base       float64
length           float64
width            float64
height           float64
curb-weight       int64
engine-type       object
num-of-cylinders  int32
engine-size       int64
fuel-system       object
bore             float64
stroke           float64
compression-ratio float64
horsepower        object
peak-rpm          float64
city-mpg          int64
highway-mpg       int64
price            float64
dtype: object

```

16. Normalize the columns "length", "width" and "height" so their value ranges from 0 to 1.

```
df['length']=(df['length']-df['length'].min())/(df['length'].max()-df['length'].min())
df['width']=(df['width']-df['width'].min())/(df['width'].max()-df['width'].min())
df['height']=(df['height']-df['height'].min())/(df['height'].max()-df['height'].min())
```

```
df.head()
```

	symboling	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location
0	3	122	alfa-romero	gas	std	2	convertible	rwd	front
1	3	122	alfa-romero	gas	std	2	convertible	rwd	front
2	1	122	alfa-romero	gas	std	2	hatchback	rwd	front
3	2	164	audi	gas	std	4	sedan	fwd	front
4	2	164	audi	gas	std	4	sedan	4wd	front

5 rows × 10 columns

Exercise 3

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from scipy.stats import norm
from scipy import stats
import warnings
warnings.filterwarnings('ignore')
```

```
data = pd.read_csv('auto-mpg.csv',index_col='car name')
data.head()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin
car name								
chevrolet chevelle malibu	18.0	8	307.0	130	3504	12.0	70	1
buick skylark 320	15.0	8	350.0	165	3693	11.5	70	1

```
print(data.head())
print(data.index)
print(data.columns)
```

	mpg	cylinders	displacement	horsepower	weight	\
car name						
chevrolet chevelle malibu	18.0	8	307.0	130	3504	
buick skylark 320	15.0	8	350.0	165	3693	
plymouth satellite	18.0	8	318.0	150	3436	
amc rebel sst	16.0	8	304.0	150	3433	
ford torino	17.0	8	302.0	140	3449	
acceleration model year origin						
car name						
chevrolet chevelle malibu	12.0	70	1			
buick skylark 320	11.5	70	1			
plymouth satellite	11.0	70	1			

```

amc rebel sst          12.0          70          1
ford torino            10.5          70          1
Index(['chevrolet chevelle malibu', 'buick skylark 320', 'plymouth satellite',
      'amc rebel sst', 'ford torino', 'ford galaxie 500', 'chevrolet impala',
      'plymouth fury iii', 'pontiac catalina', 'amc ambassador dpl',
      ...,
      'chrysler lebaron medallion', 'ford granada l', 'toyota celica gt',
      'dodge charger 2.2', 'chevrolet camaro', 'ford mustang gl', 'vw pickup',
      'dodge rampage', 'ford ranger', 'chevy s-10'],
      dtype='object', name='car name', length=398)
Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
      'acceleration', 'model year', 'origin'],
      dtype='object')

```

```
data.shape
```

```
(398, 8)
```

```
data.isnull().any()
```

```

mpg             False
cylinders       False
displacement    False
horsepower      False
weight          False
acceleration    False
model year      False
origin          False
dtype: bool

```

```
data.dtypes
```

```

mpg             float64
cylinders       int64
displacement    float64
horsepower      object
weight          int64
acceleration    float64
model year      int64
origin          int64
dtype: object

```

```
data.horsepower.unique()
```

```

array(['130', '165', '150', '140', '198', '220', '215', '225', '190',
      '170', '160', '95', '97', '85', '88', '46', '87', '90', '113',
      '200', '210', '193', '?', '100', '105', '175', '153', '180', '110',
      '72', '86', '70', '76', '65', '69', '60', '80', '54', '208', '155',
      '112', '92', '145', '137', '158', '167', '94', '107', '230', '49',
      '75', '91', '122', '67', '83', '78', '52', '61', '93', '148',
      '129', '96', '71', '98', '115', '53', '81', '79', '120', '152',
      '102', '108', '68', '58', '149', '89', '63', '48', '66', '139',
      '103', '125', '133', '138', '135', '142', '77', '62', '132', '84',
      '64', '74', '116', '82'], dtype=object)

```

```
data = data[data.horsepower != '?']
```

```
print('? ' in data.horsepower)
```

```
False
```

```
data.shape
```

```
(392, 8)
```

```
data.dtypes
```

```

mpg             float64
cylinders       int64

```

```
displacement    float64
horsepower      object
weight          int64
acceleration     float64
model_year      int64
origin          int64
dtype: object

data.horsepower = data.horsepower.astype('float')
data.dtypes

mpg             float64
cylinders       int64
displacement    float64
horsepower      float64
weight          int64
acceleration     float64
model_year      int64
origin          int64
dtype: object

data.describe()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	origin
count	392.000000	392.000000	392.000000	392.000000	392.000000	392.000000	392.000000
mean	23.445918	5.471939	194.411990	104.469388	2977.584184	15.541327	75.973958
std	7.805007	1.705783	104.644004	38.491160	849.402560	2.758864	3.658812
min	9.000000	3.000000	68.000000	46.000000	1613.000000	8.000000	70.000000
25%	17.000000	4.000000	105.000000	75.000000	2225.250000	13.775000	73.000000
50%	22.750000	4.000000	151.000000	93.500000	2803.500000	15.500000	76.000000
75%	29.000000	8.000000	275.750000	126.000000	3614.750000	17.025000	79.000000

```
data.mpg.describe()

count    392.000000
mean     23.445918
std       7.805007
min       9.000000
25%      17.000000
50%      22.750000
75%      29.000000
max      46.600000
Name: mpg, dtype: float64

sns.distplot(data['mpg'])
```



```
<Axes: xlabel='mpg', ylabel='Density'>
```



```
print("Skewness: %f" % data['mpg'].skew())
print("Kurtosis: %f" % data['mpg'].kurt())
```

```
Skewness: 0.457092
Kurtosis: -0.515993
```

```
-
```



```
def scale(a):
    b = (a-a.min())/(a.max()-a.min())
    return b
```

```
data_scale = data.copy()
```

```
data_scale['displacement'] = scale(data_scale['displacement'])
data_scale['horsepower'] = scale(data_scale['horsepower'])
data_scale['acceleration'] = scale(data_scale['acceleration'])
data_scale['weight'] = scale(data_scale['weight'])
data_scale['mpg'] = scale(data_scale['mpg'])
```

```
data_scale.head()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	c
car name								
chevrolet chevelle malibu	0.239362	8	0.617571	0.456522	0.536150	0.238095	70	
buick skylark 320	0.159574	8	0.728682	0.646739	0.589736	0.208333	70	

```
data['Country_code'] = data.origin.replace([1,2,3],['USA','Europe','Japan'])
data_scale['Country_code'] = data.origin.replace([1,2,3],['USA','Europe','Japan'])
```

```
data_scale.head()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	c
car name								
chevrolet chevelle malibu	0.239362	8	0.617571	0.456522	0.536150	0.238095	70	
buick skylark 320	0.159574	8	0.728682	0.646739	0.589736	0.208333	70	


```
var = 'Country_code'
data_plt = pd.concat([data_scale['mpg'], data_scale[var]], axis=1)
f, ax = plt.subplots(figsize=(8, 6))
fig = sns.boxplot(x=var, y="mpg", data=data_plt)
fig.axis(ymin=0, ymax=1)
plt.axhline(data_scale.mpg.mean(),color='r',linestyle='dashed',linewidth=2)
```

```
<matplotlib.lines.Line2D at 0x13e346c5f40>
```




```
var = 'model year'
data_plt = pd.concat([data_scale['mpg'], data_scale[var]], axis=1)
f, ax = plt.subplots(figsize=(8, 6))
fig = sns.boxplot(x=var, y="mpg", data=data_plt)
fig.axis(ymin=0, ymax=1)
plt.axhline(data_scale.mpg.mean(),color='r',linestyle='dashed',linewidth=2)
```

```
<matplotlib.lines.Line2D at 0x13e37ac0980>
```



```
var = 'cylinders'
data_plt = pd.concat([data_scale['mpg'], data_scale[var]], axis=1)
f, ax = plt.subplots(figsize=(8, 6))
fig = sns.boxplot(x=var, y="mpg", data=data_plt)
fig.axis(ymin=0, ymax=1)
plt.axhline(data_scale.mpg.mean(),color='r',linestyle='dashed',linewidth=2)
```

<matplotlib.lines.Line2D at 0x13e37fbb290>



```
corrmat = data.corr()
f, ax = plt.subplots(figsize=(12, 9))
sns.heatmap(corrmat, square=True)
```

```
-----
ValueError                                Traceback (most recent call last)
c:\Users\aadit\Desktop\BTech\S5\Foundations Of Data Science\LAB\lab 2\lab2.ipynb Cell
57 line 1
----> <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20of%20Data%20Science/LAB/lab%202/lab
line=0'>1</a> corrmat = data.corr()
      <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20of%20Data%20Science/LAB/lab%202/lab
line=1'>2</a> f, ax = plt.subplots(figsize=(12, 9))
      <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20of%20Data%20Science/LAB/lab%202/lab
line=2'>3</a> sns.heatmap(corrmat, square=True)
```

```
File ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\frame.py:10704, in
DataFrame.corr(self, method, min_periods, numeric_only)
    10702 cols = data.columns
    10703 idx = cols.copy()
-> 10704 mat = data.to_numpy(dtype=float, na_value=np.nan, copy=False)
    10706 if method == "pearson":
    10707     correl = libalgos.nancorr(mat, minp=min_periods)
```


```
File ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\frame.py:1889, in
DataFrame.to_numpy(self, dtype, copy, na_value)
    1887 if dtype is not None:
    1888     dtype = np.dtype(dtype)
-> 1889 result = self._mgr.as_array(dtype=dtype, copy=copy, na_value=na_value)
    1890 if result.dtype is not dtype:
    1891     result = np.array(result, dtype=dtype, copy=False)
```

```
File ~\AppData\Roaming\Python\Python312\site-
packages\pandas\core\internals\managers.py:1656, in BlockManager.as_array(self, dtype,
copy, na_value)
    1654     arr.flags.writeable = False
    1655 else:
-> 1656     arr = self._interleave(dtype=dtype, na_value=na_value)
    1657     # The underlying data was copied within _interleave, so no need
    1658     # to further copy if copy=True or setting na_value
    1660 if na_value is lib.no_default:
```

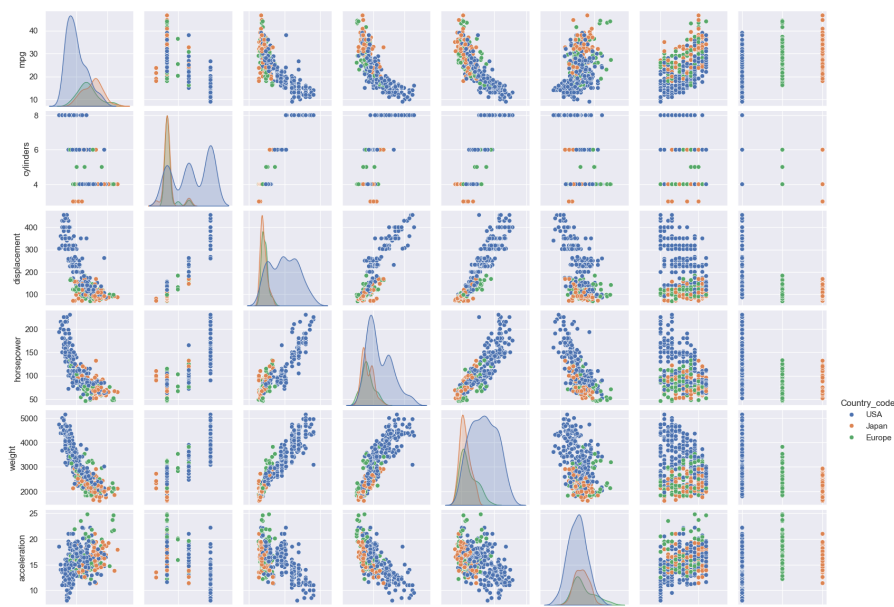
```
File ~\AppData\Roaming\Python\Python312\site-
packages\pandas\core\internals\managers.py:1715, in BlockManager._interleave(self,
dtype, na_value)
```

```
factors = ['cylinders', 'displacement', 'horsepower', 'acceleration', 'weight', 'mpg']
corrmat = data[factors].corr()
f, ax = plt.subplots(figsize=(12, 9))
sns.heatmap(corrmat, square=True)
```

<Axes: >



```
sns.set()  
sns.pairplot(data, size = 2.0, hue = 'Country_code')  
plt.show()
```



data.index

```
Index(['chevrolet chevelle malibu', 'buick skylark 320', 'plymouth satellite',
      'amc rebel sst', 'ford torino', 'ford galaxie 500', 'chevrolet impala',
      'plymouth fury iii', 'pontiac catalina', 'amc ambassador dpl',
      ...,
      'chrysler lebaron medallion', 'ford granada l', 'toyota celica gt',
      'dodge charger 2.2', 'chevrolet camaro', 'ford mustang gl', 'vw pickup',
      'dodge rampage', 'ford ranger', 'chevy s-10'],
      dtype='object', name='car name', length=392)
```

```
data[data.index.str.contains('subaru')].index.str.replace('(.*)', 'subaru dl')
```


```
Index(['subaru', 'subaru dl', 'subaru dl', 'subaru'], dtype='object', name='car name')
```

```
data['Company_Name'] = data.index.str.extract('^.*?\s')
```

```
data['Company_Name'] = data['Company_Name'].replace(['volkswagen', 'vokswagen', 'vw'], 'VW')
data['Company_Name'] = data['Company_Name'].replace('maxda', 'mazda')
data['Company_Name'] = data['Company_Name'].replace('toyouta', 'toyota')
data['Company_Name'] = data['Company_Name'].replace('mercedes', 'mercedes-benz')
data['Company_Name'] = data['Company_Name'].replace('nissan', 'datsun')
data['Company_Name'] = data['Company_Name'].replace('capri', 'ford')
data['Company_Name'] = data['Company_Name'].replace(['chevroelt', 'chevy'], 'chevrolet')
data['Company_Name'].fillna(value = 'subaru', inplace=True)
```

```
var = 'Company_Name'
data_plt = pd.concat([data_scale['mpg'], data[var]], axis=1)
f, ax = plt.subplots(figsize=(20,10))
fig = sns.boxplot(x=var, y="mpg", data=data_plt)
fig.set_xticklabels(ax.get_xticklabels(), rotation=30)
fig.axis(ymin=0, ymax=1)
plt.axhline(data_scale.mpg.mean(), color='r', linestyle='dashed', linewidth=2)
```

```
<matplotlib.lines.Line2D at 0x13e3a021e20>
```



```
data.Company_Name.isnull().any()
```

```
False
```

```
var='mpg'
data[data[var]== data[var].min()]
```

```
      mpg  cylinders  displacement  horsepower  weight  acceleration  model  origin  Co
      year
car
name
```

```
data[data[var]== data[var].max()]
```

```
      mpg  cylinders  displacement  horsepower  weight  acceleration  model  origin  C
      year
car
name
```

```
var='displacement'
data[data[var]== data[var].min()]
```

```
      mpg  cylinders  displacement  horsepower  weight  acceleration  model  origin  Cou
      year
car
name
```

```
data[data[var]== data[var].max()]
```

```
      mpg  cylinders  displacement  horsepower  weight  acceleration  model  origin
      year
car
name
```

car name	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin
pontiac catalina	14.0	8	455.0	225.0	4425	10.0	70	1

```
var='horsepower'
data[data[var]== data[var].min()]
```

```

    mpg cylinders displacement horsepower weight acceleration model
                                year origin
car name
volkswagen
data[data[var]== data[var].max()]

```

```

    mpg cylinders displacement horsepower weight acceleration model
                                year origin
car name

```

```

var='weight'
data[data[var]== data[var].min()]

```

```

    mpg cylinders displacement horsepower weight acceleration model
                                year origin
car name

```

```

data[data[var]== data[var].max()]

```

```

    mpg cylinders displacement horsepower weight acceleration model
                                year origin
car name

```

```

var='acceleration'
data[data[var]== data[var].min()]

```

```

    mpg cylinders displacement horsepower weight acceleration model
                                year origin
car name

```

```

data[data[var]== data[var].max()]

```

```

    mpg cylinders displacement horsepower weight acceleration model
                                year origin
car name

```

```

var = 'horsepower'
plot = sns.lmplot(var, 'mpg', data=data, hue='Country_code')
plot.set(ylim = (0,50))

```

```

-----
TypeError                                Traceback (most recent call last)
c:\Users\aadit\Desktop\BTech\S5\Foundations Of Data Science\LAB\lab 2\lab2.ipynb Cell
76 line 2
      <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20Of%20Data%20Science/LAB/lab%202/lab
line=0'>1</a> var = 'horsepower'
----> <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20Of%20Data%20Science/LAB/lab%202/lab
line=1'>2</a> plot = sns.lmplot(var, 'mpg', data=data, hue='Country_code')
      <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20Of%20Data%20Science/LAB/lab%202/lab

```

```

var = 'displacement'
plot = sns.lmplot(var, 'mpg', data=data, hue='Country_code')
plot.set(ylim = (0,50))

```

```

-----
TypeError                                Traceback (most recent call last)
c:\Users\aadit\Desktop\BTech\S5\Foundations Of Data Science\LAB\lab 2\lab2.ipynb Cell
77 line 2
      <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20Of%20Data%20Science/LAB/lab%202/lab2.ipynb?line=0'>1</a> var = 'displacement'
      <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations Of Data Science\LAB\lab 2\lab2.ipynb Cell
77 line 2
var = 'weight'
plot = sns.lmplot(var, 'mpg', data=data, hue='Country_code')
plot.set(ylim = (0,50))

```

```

-----
TypeError                                Traceback (most recent call last)
c:\Users\aadit\Desktop\BTech\S5\Foundations Of Data Science\LAB\lab 2\lab2.ipynb Cell
78 line 2
      <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20Of%20Data%20Science/LAB/lab%202/lab2.ipynb?line=0'>1</a> var = 'weight'
----> <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20Of%20Data%20Science/LAB/lab%202/lab2.ipynb?line=1'>2</a> plot = sns.lmplot(var, 'mpg', data=data, hue='Country_code')
      <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20Of%20Data%20Science/LAB/lab%202/lab2.ipynb?line=1'>2</a> plot = sns.lmplot(var, 'mpg', data=data, hue='Country_code')

```

```

var = 'acceleration'
plot = sns.lmplot(var, 'mpg', data=data, hue='Country_code')
plot.set(ylim = (0,50))

```

```

-----
TypeError                                Traceback (most recent call last)
c:\Users\aadit\Desktop\BTech\S5\Foundations Of Data Science\LAB\lab 2\lab2.ipynb Cell
79 line 2
      <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20Of%20Data%20Science/LAB/lab%202/lab2.ipynb?line=0'>1</a> var = 'acceleration'
----> <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20Of%20Data%20Science/LAB/lab%202/lab2.ipynb?line=1'>2</a> plot = sns.lmplot(var, 'mpg', data=data, hue='Country_code')
      <a href='vscode-notebook-
cell:/c%3A/Users/aadit/Desktop/BTech/S5/Foundations%20Of%20Data%20Science/LAB/lab%202/lab2.ipynb?line=1'>2</a> plot = sns.lmplot(var, 'mpg', data=data, hue='Country_code')

```

```
data['Power_to_weight'] = ((data.horsepower*0.7457)/data.weight)
```

```
data.sort_values(by='Power_to_weight', ascending=False ).head()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin
car name								
buick estate wagon (sw)	14.0	8	455.0	225.0	3086	10.0	70	1
pontiac grand seal	16.0	8	400.0	230.0	4278	9.5	73	1

```
data.head()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin
car name								
chevrolet chevelle malibu	18.0	8	307.0	130.0	3504	12.0	70	1
buick skylark 320	15.0	8	350.0	165.0	3693	11.5	70	1


```

from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from math import sqrt
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import KFold

```

```

factors = ['cylinders', 'displacement', 'horsepower', 'acceleration', 'weight', 'origin', 'model year']
X = pd.DataFrame(data[factors].copy())
y = data['mpg'].copy()

```

```
X = StandardScaler().fit_transform(X)
```

```

X_train,X_test,y_train,y_test=train_test_split(X,y,test_size = 0.33,random_state=324)
X_train.shape[0] == y_train.shape[0]

```

```
True
```

```
regressor = LinearRegression()
```

```
regressor.get_params()
```

```
{'copy_X': True, 'fit_intercept': True, 'n_jobs': None, 'positive': False}
```

```
regressor.fit(X_train,y_train)
```

```

LinearRegression
LinearRegression()

```

```
y_predicted = regressor.predict(X_test)
```

```

rmse = sqrt(mean_squared_error(y_true=y_test,y_pred=y_predicted))
rmse

```

```
3.486729614901561
```

```

gb_regressor = GradientBoostingRegressor(n_estimators=4000)
gb_regressor.fit(X_train,y_train)

```

```

GradientBoostingRegressor
GradientBoostingRegressor(n_estimators=4000)

```

```
gb_regressor.get_params()
```

```

{'alpha': 0.9,
 'ccp_alpha': 0.0,
 'criterion': 'friedman_mse',
 'init': None,
 'learning_rate': 0.1,
 'loss': 'squared_error',
 'max_depth': 3,
 'max_features': None,
 'max_leaf_nodes': None,
 'min_impurity_decrease': 0.0,
 'min_samples_leaf': 1,
 'min_samples_split': 2,
 'min_weight_fraction_leaf': 0.0,
 'n_estimators': 4000,
 'n_iter_no_change': None,
 'random_state': None,
 'subsample': 1.0,
 'tol': 0.0001,
 'validation_fraction': 0.1,
 'verbose': 0,
 'warm_start': False}

```


```
y_predicted_gbr = gb_regressor.predict(X_test)
```

```
rmse_bgr = sqrt(mean_squared_error(y_true=y_test,y_pred=y_predicted_gbr))
rmse_bgr
```

```
2.678422792051918
```

```
fi= pd.Series(gb_regressor.feature_importances_,index=factors)
fi.plot.barh()
```

```
<Axes: >
```



```
from sklearn.decomposition import PCA
pca = PCA(n_components=2)
pca.fit(data[factors])
```

```
PCA
PCA(n_components=2)
```

```
pca.explained_variance_ratio_
array([0.99756151, 0.0020628 ])
```

```
pca1 = pca.components_[0]
pca2 = pca.components_[1]
```

```
transformed_data = pca.transform(data[factors])
pc1 = transformed_data[:,0]
pc2 = transformed_data[:,1]
plt.scatter(pc1,pc2)
```

```
<matplotlib.collections.PathCollection at 0x13e32fec410>
c = pca.inverse_transform(transformed_data[(transformed_data[:,0]>0 )& (transformed_data[:,1]>250)])
factors
```

```
['cylinders',
 'displacement',
 'horsepower',
 'acceleration',
 'weight',
 'origin',
 'model year']
```



```
c
array([[ 9.32016159e+00,  4.65727261e+02,  1.90441442e+02,  5.95699243e+00,
        3.08611199e+03,  6.23550659e-01,  6.93571097e+01]])
-50
data[(data['model year'] == 70 )&( data.displacement>400)]
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	Country_code	Company_Name	Power_to_weight
car name											
ford galaxie 500	15.0	8	429.0	198.0	4341	10.0	70	1	USA	subaru	0.034013
chevrolet impala	14.0	8	454.0	220.0	4354	9.0	70	1	USA	subaru	0.037679
plymouth fury iii	14.0	8	440.0	215.0	4312	8.5	70	1	USA	subaru	0.037181

```
cv_sets = KFold(n_splits=10, shuffle= True,random_state=100)
params = {'n_estimators' : list(range(40,61)),
          'max_depth' : list(range(1,10)),
          'learning_rate' : [0.1,0.2,0.3] }
grid = GridSearchCV(gb_regressor, params,cv=cv_sets,n_jobs=4)
grid = grid.fit(X_train, y_train)
grid.best_estimator_
```

▼

GradientBoostingRegressor

GradientBoostingRegressor(learning_rate=0.3, max_depth=2, n_estimators=55)

```
gb_regressor_t = grid.best_estimator_
gb_regressor_t.fit(X_train,y_train)
```

▼

GradientBoostingRegressor

GradientBoostingRegressor(learning_rate=0.3, max_depth=2, n_estimators=55)

```
y_predicted_gbr_t = gb_regressor_t.predict(X_test)
rmse = sqrt(mean_squared_error(y_true=y_test,y_pred=y_predicted_gbr_t))
rmse
```

2.6762648633586865

```
data.duplicated().any()
False
```